

You have an Azure subscription that contains an Azure Storage account.

You plan to create an Azure container instance named container1 that will use a Docker image named Image1. Image1 contains a Microsoft SQL Server instance that requires persistent storage.

You need to configure a storage service for Container1.

What should you use?

Azure Files

Azure Blob storage

Azure Queue storage

Azure Table storage

You have an app named App1 that runs on two Azure virtual machines named VM1 and VM2.

You plan to implement an Azure Availability Set for App1. The solution must ensure that App1 is available during planned maintenance of the hardware hosting VM1 and VM2.

What should you include in the Availability Set?

- one update domain
- two fault domains
- one fault domain
- two update domains

You have an on-premises server that contains a folder named D:\Folder1.

You need to copy the contents of D:\Folder1 to the public container in an Azure Storage account named contosodata.

Which command should you run?

- . https://contosodata.blob.core.windows.net/public
- azcopy sync D:\folder1 https://contosodata.blob.core.windows.net/public --snapshot
- azcopy copy D:\folder1 https://contosodata.blob.core.windows.net/public --recursive
- az storage blob copy start-batch D:\Folder1 https://contosodata.blob.core.windows.net/public

You have an Azure subscription named Subscription1 that contains the storage accounts shown in the following table:

| Name | Account Kind | Azure service that contains data |
|----------|--------------------------------|----------------------------------|
| storage1 | Storage | File |
| storage2 | StorageV2 (general purpose v2) | File, Table |
| storage3 | StorageV2 (general purpose v2) | Queue |
| storage4 | BlobStorage | Blob |

You plan to use the Azure Import/Export service to export data from Subscription1.

You need to identify which storage account can be used to export the data.

What should you identify

storage1

storage2

storage3

storage4

You have an Azure subscription that contains the storage accounts shown in the following table.

| Name | Kind | Performance | Replication | Access tier |
|----------|--------------------------------|-------------|--|-------------|
| storage1 | Storage (general purpose v1) | Premium | Geo-redundant storage (GRS) | None |
| storage2 | StorageV2 (general purpose v2) | Standard | Locally-redundant storage (LRS) | Cool |
| storage3 | StorageV2 (general purpose v2) | Premium | Read-access georedundant storage(RA-GRS) | Hot |
| storage4 | BlobStorage | Standard | Locally-redundant storage (LRS) | Hot |

You need to identify which storage account can be converted to zone-redundant storage (ZRS) replication by requesting a live migration from Azure support.

What should you identify?

storage1

storage2

storage3

storage4

You have an Azure subscription that contains a storage account named account1.

You plan to upload the disk files of a virtual machine to account1 from your on-premises network. The on-premises network uses a public IP address space of 131.107.1.0/24.

You plan to use the disk files to provision an Azure virtual machine named VM1. VM1 will be attached to a virtual network named VNet1. VNet1 uses an IP address space of 192.168.0.0/24.

You need to configure account1 to meet the following requirements:

- ⇒ Ensure that you can upload the disk files to account1.
- ⇒ Ensure that you can attach the disks to VM1.
- ⇒ Prevent all other access to account1.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- From the Networking blade of account1, select Selected networks.
- From the Networking blade of account1, select Allow trusted Microsoft services to access this storage account.
- From the Networking blade of account1, add the 131.107.1.0/24 IP address range.
- From the Networking blade of account1, add VNet1.
- From the Service endpoints blade of VNet1, add a service endpoint

DRAG DROP

You have an on-premises file server named Server1 that runs Windows Server 2016.

You have an Azure subscription that contains an Azure file share.

You deploy an Azure File Sync Storage Sync Service, and you create a sync group.

You need to synchronize files from Server1 to Azure.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Action

1. Install the Azure File Sync agent on Server1
2. Create an Azure on-premises data gateway
3. Create a Recovery Services vault
4. Register Server1
5. Add a server endpoint
6. Install the DFS Replication server role on Server1

1,4,5

5,3,2

6,3,5

1,3,5

HOTSPOT

You plan to create an Azure Storage account in the Azure region of East US 2.

You need to create a storage account that meets the following requirements:

- ✚ Replicates synchronously.
- ✚ Remains available if a single data center in the region fails.

How should you configure the storage account? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Replication: _____

- Geo-redundant storage (GRS)
- Locally-redundant storage (LRS)
- Read-access geo-redundant storage (RA GRS)
- Zone-redundant storage (ZRS)

Account type: _____

- Blob storage
- Storage (general purpose v1)
- StorageV2 (general purpose v2)

HOTSPOT

You have an Azure Storage account named storage1 that uses Azure Blob storage and Azure File storage.

You need to use AzCopy to copy data to the blob storage and file storage in storage1.

Which authentication method should you use for each type of storage?

To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Blob storage: _____

- Azure Active Directory (Azure AD) only
- Shared access signatures (SAS) only
- Access keys and shared access signatures (SAS) only
- Azure Active Directory (Azure AD) and shared access signatures (SAS) only
- Azure Active Directory (Azure AD), access keys, and shared access signatures (SAS)

File storage: _____

- Azure Active Directory (Azure AD) only
- Shared access signatures (SAS) only
- Access keys and shared access signatures (SAS) only
- Azure Active Directory (Azure AD) and shared access signatures (SAS) only
- Azure Active Directory (Azure AD), access keys, and shared access signatures (SAS)

HOTSPOT

You have an Azure subscription named Subscription1 that contains the resources shown in the following table:

| Name | Type | Location | Resource group |
|------------|-------------------------|-------------|----------------|
| RG1 | Resource group | East US | Not applicable |
| RG2 | Resource group | West US | Not applicable |
| Vault1 | Recovery Services vault | West Europe | RG1 |
| storage1 | Storage account | East US | RG2 |
| storage2 | Storage account | West US | RG1 |
| storage3 | Storage account | West Europe | RG2 |
| Analytics1 | Log Analytics workspace | East US | RG1 |
| Analytics2 | Log Analytics workspace | West US | RG2 |
| Analytics3 | Log Analytics workspace | West Europe | RG1 |

You plan to configure Azure Backup reports for Vault1.

You are configuring the Diagnostics settings for the AzureBackupReports log.

Which storage accounts and which Log Analytics workspaces can you use for the Azure Backup reports of Vault1? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Storage accounts: _____

- storage1 only
- storage 2 only
- storage3 only
- storage1, storage 2 , and storage 3

Log Analytics workspaces: _____

- Analytics1 only
- Analytics2 only
- Analytics3 only
- Analytics1, Analytics2, and Analytics3

HOTSPOT

You have an Azure subscription named Subscription1.

In Subscription1, you create an Azure file share named share1.

You create a shared access signature (SAS) named SAS1 as shown in the following exhibit:

Allowed services ●

Blob File Queue Table

Allowed resource types ●

Service Container Object

Allowed permissions ●

Read Write Delete List Add Create Update Process

Start and expiry date/time ●

Start
2018-09-01 2:00:00 PM

End
2018-09-14 2:00:00 PM

(UTC+02:00) --- Current Timezone ---

Allowed IP addresses ●

193.77.134.10-193.77.134.50

Allowed protocols ●

HTTPS only HTTPS and HTTP

Signing key ●

key1

Generate SAS and connection string

To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

If on September 2, 2018, you run Microsoft Azure Storage Explorer on a computer that has an IP address of 193.77.134.1, and you use SAS1 to connect to the storage account, you [answer choice]. _____

- will be prompted for credentials
- will have no access
- will have read, write, and list access
- will have read-only access

If on September 10, 2018, you run the net use command on a computer that has an IP address of 193.77.134.50, and you use SAS1 as the password to connect to share1, you [answer choice]. _____

- will be prompted for credentials
- will have no access
- will have read, write, and list access
- Will have read-only access

HOTSPOT You have Azure subscription that includes following Azure file shares:

| name | In storage account | Location |
|--------|--------------------|----------|
| share1 | storage1 | West US |
| share2 | storage1 | West US |

You have the following on-premises servers:

| Name | Folders |
|---------|------------------------|
| Server1 | D:-Folder1, E:YFolder2 |
| Server2 | D:-Data |

You create a Storage Sync Service named Sync1 and an Azure File Sync group named Group1. Group1 uses share1 as a cloud endpoint.

You register Server1 and Server2 in Sync1. You add D:\Folder1 on Server1 as a server endpoint of Group1.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

share 2 can be added as a cloud endpoint for Group 1 _____

Yes

No

E:\Folder2 on Server1 can be added as a server endpoint for Group1 _____

Yes

No

D:\Data on Server2 can be added as a server endpoint for Group1 _____

Yes

No

DRAG DROP

You have an Azure subscription named Subscription1.

You create an Azure Storage account named contosostorage, and then you create a file share named data.

Which UNC path should you include in a script that references files from the data file share? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once or not at all. You may need to drag

the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

| Values | |
|------------------|-----------------------|
| blob | blob.core.windows.net |
| contosostorage | data |
| file | file.core.windows.net |
| portal.azure.com | subscription1 |

Answer Area

W \ \

contosostorage

file.core.windows.net

data

blob

subscription1

portal.azure.com

blob.core.windows.net

file

data

file.core. Windows net

file

contosostorage

HOTSPOT

You have an Azure subscription that contains an Azure Storage account.

You plan to copy an on-premises virtual machine image to a container named `vmimages`.

You need to create the container for the planned image.

Which command should you run? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

azcopy _____

make

sync

copy

blob

dfs

queue

table

images

file

HOTSPOT

You have an Azure File sync group that has the endpoints shown in the following table.

| Name | Type |
|-----------|-----------------|
| Endpoint1 | Cloud endpoint |
| Endpoint2 | Server endpoint |
| Endpoint3 | Server endpoint |

Cloud tiering is enabled for Endpoint3.

You add a file named File1 to Endpoint1 and a file named File2 to Endpoint2.

On which endpoints will File1 and File2 be available within 24 hours of adding the files? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

File1: _____

- Endpoint1 only
- Endpoint3 only
- Endpoint2 and Endpoint3 only
- Endpoint1, Endpoint2, and Endpoint3

File2 _____

- Endpoint2 only
- Endpoint 3 only
- Endpoint2 and Endpoint 3 only
- Endpoint1, Endpoint2, and Endpoint3

HOTSPOT

You have several Azure virtual machines on a virtual network named VNet1.

You configure an Azure Storage account as shown in the following exhibit.

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

The virtual machines on the 10.2.9.0/24 subnet will have network connectivity to the fil shares in the storage account _____

- always
- during a backup
- never

Azure Backup will be able to back up the unmanaged hard disks of the virtual machines in the storage account [answer choice]. _____

- always
- during a backup
- never

Validate ✓

Solution:

Explanation:

Box 1: Never

Box 2: Never

After you configure firewall and virtual network settings for your storage account, select Allow trusted Microsoft services to access this storage account as an exception to enable Azure Backup service to access the network restricted storage account.

Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-how-to-use-files-windows>

<https://azure.microsoft.com/en-us/blog/azure-backup-now-supports-storage-accounts-secured-with-azure-storage-firewalls-and-virtual-networks/>

HOTSPOT

You have a sync group named Sync1 that has a cloud endpoint. The cloud endpoint includes a file named File1.txt.

Your on-premises network contains servers that run Windows Server 2016. The servers are configured as shown in the following table.

| Name | Share | Share contents |
|---------|--------|----------------------|
| Server1 | Share1 | File1.txt, File2.txt |
| Server2 | Share2 | File2.txt, File3.txt |

You add Share1 as an endpoint for Sync1. One hour later, you add Share2 as an endpoint for Sync1.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

On the cloud endpoint, File1.txt is overwritten by File1.txt from Share1. _____

Yes

No

On Server1, File1.txt is overwritten by File1.txt from the cloud endpoint. _____

Yes

No

File1.txt from Share1 replicates to Share2. _____

Yes

No

You have Azure Storage accounts as shown in the following exhibit

| Storage accounts | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|-----------------|----------------|------------|--------------|----------------|----------------------|------|------|----------------|----------|--------------|-------------|-------------|-----------------|-----------------|---------|------------|---------|----------------|-------------------|-----------------|-----------------|-----------|------------|------------|----------------|------------------|-----------------|-----------------|-------------|------------|---------|----------------|----------------------|
| ✖ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| + Add Edit columns Refresh Assign Tags Delete | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| Subscription: All 2 selected - Don't see a subscription? Switch directories | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| Filter by name... All subscriptions All resource groups All types All locations No grouping | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 items | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| <table border="1"> <thead> <tr> <th>NAME</th> <th>TYPE</th> <th>RESOURCE GROUP</th> <th>LOCATION</th> <th>SUBSCRIPTION</th> <th>ACCESS TIER</th> <th>REPLICAT...</th> </tr> </thead> <tbody> <tr> <td>storageaccount1</td> <td>Storage account</td> <td>Storage</td> <td>ContosoRG1</td> <td>East US</td> <td>Subscription 1</td> <td>Read-access ge...</td> </tr> <tr> <td>storageaccount2</td> <td>Storage account</td> <td>StorageV2</td> <td>ContosoRG1</td> <td>Central US</td> <td>Subscription 1</td> <td>Geo-redundant...</td> </tr> <tr> <td>storageaccount3</td> <td>Storage account</td> <td>BlobStorage</td> <td>ContosoRG1</td> <td>East US</td> <td>Subscription 1</td> <td>Locally-redundant...</td> </tr> </tbody> </table> | | | | | | | NAME | TYPE | RESOURCE GROUP | LOCATION | SUBSCRIPTION | ACCESS TIER | REPLICAT... | storageaccount1 | Storage account | Storage | ContosoRG1 | East US | Subscription 1 | Read-access ge... | storageaccount2 | Storage account | StorageV2 | ContosoRG1 | Central US | Subscription 1 | Geo-redundant... | storageaccount3 | Storage account | BlobStorage | ContosoRG1 | East US | Subscription 1 | Locally-redundant... |
| NAME | TYPE | RESOURCE GROUP | LOCATION | SUBSCRIPTION | ACCESS TIER | REPLICAT... | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| storageaccount1 | Storage account | Storage | ContosoRG1 | East US | Subscription 1 | Read-access ge... | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| storageaccount2 | Storage account | StorageV2 | ContosoRG1 | Central US | Subscription 1 | Geo-redundant... | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| storageaccount3 | Storage account | BlobStorage | ContosoRG1 | East US | Subscription 1 | Locally-redundant... | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

You can use [answer choice] for Azure Table Storage. _____

- storageaccount1 only
- storageaccount2 only
- storageaccount3 only
- storageaccount 1 and storageaccount2 only
- storageaccount2 and storageaccount 3 only

You can use [answer choice] for Azure Blob storage _____

- storageaccount 3 only
- storageaccount2 and storageaccount 3 only
- storageaccount1 and storageaccount 3 only
- all the storage accounts

Validate ✓

Solution:

Explanation:

Box 1: storageaccount1 and storageaccount2 only

Box 2: All the storage accounts

Note: The three different storage account options are: General-purpose v2 (GPv2) accounts, General-purpose v1 (GPv1) accounts, and Blob storage accounts.

* General-purpose v2 (GPv2) accounts are storage accounts that support all of the latest features for blobs, files, queues, and tables.

* Blob storage accounts support all the same block blob features as GPv2, but are limited to supporting only block blobs.

* General-purpose v1 (GPv1) accounts provide access to all Azure Storage services, but may not have the latest features or the lowest per gigabyte pricing.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-options>

You have Azure subscription that includes data in following locations:

| Name | Type |
|------------|-------------------|
| container1 | Blob container |
| share1 | Azure files share |
| DB1 | SQL database |
| Table1 | Azure Table |

You plan to export data by using Azure import/export job named Export1.

You need to identify the data that can be exported by using Export1.

Which data should you identify?

DB1

container1

Share1

Table1

HOTSPOT

You have an Azure Storage account named storage1.

You have an Azure App Service app named App1 and an app named App2 that runs in an Azure container instance. Each app uses a managed identity.

You need to ensure that App1 and App2 can read blobs from storage1. The solution must meet the following requirements:

- ◆ Minimize the number of secrets used.
- ◆ Ensure that App2 can only read from storage1 for the next 30 days.

What should you configure in storage1 for each app? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

App1 _____

- Access keys
- Advanced security
- Access control (IAM)
- Shared access signatures (SAS)

Blank(i) _____

- Access keys
- Advanced security
- Access control (IAM)
- Shared access signatures (SAS)

HOTSPOT

You need to create an Azure Storage account that meets the following requirements:

- Minimizes costs
- Supports hot, cool, and archive blob tiers
- Provides fault tolerance if a disaster affects the Azure region where the account resides

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

az storage account create -g RG1 -n storageaccount1 --kind _____

- BlobStorage
- Storage
- StorageV2

az storage account create -g RG1 -n storageaccount1 --sku _____

- Standard_GRS
- Standard_LRS
- Standard_RAGRS
- Premium_LRS

Validate ✓

Solution:**Explanation:**

Box 1: StorageV2 You may only tier your object storage data to hot, cool, or archive in Blob storage and General Purpose v2 (GPv2) accounts. General Purpose v1 (GPv1) accounts do not support tiering.

General-purpose v2 accounts deliver the lowest per-gigabyte capacity prices for Azure Storage, as well as industry-competitive transaction prices.

Box 2: Standard_GRS

Geo-redundant storage (GRS): Cross-regional replication to protect against region-wide unavailability.

Incorrect Answers:

Locally-redundant storage (LRS): A simple, low-cost replication strategy. Data is replicated within a single storage scale unit.

Read-access geo-redundant storage (RA-GRS): Cross-regional replication with read access to the replica. RA-GRS provides read-only access to the data in the secondary location, in addition to geo-replication across two regions, but is

more expensive compared to GRS.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy-grs>

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers>

You have an Azure subscription that contains the resources in the following table

| Name | Type |
|--------|-----------------------|
| RG1 | Resource group |
| store1 | Azure Storage account |
| Sync1 | Azure File Sync |

Store1 contains a file share named data. Data contains 5,000 files.

You need to synchronize the files in the file share named data to an on-premises server named Server1.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

Create a container instance

Register Server1

Install the Azure File Sync agent on Server1

Download an automation script

Create a sync group

Validate ✓

Solution:

Explanation:

Step 1 (C): Install the Azure File Sync agent on Server1

The Azure File Sync agent is a downloadable package that enables Windows Server to be synced with an Azure file share

Step 2 (B): Register Server1.

Register Windows Server with Storage Sync Service

Registering your Windows Server with a Storage Sync Service establishes a trust relationship between your server (or cluster) and the Storage Sync Service.

Step 3 (E): Create a sync group and a cloud endpoint.

A sync group defines the sync topology for a set of files. Endpoints within a sync group are kept in sync with each other. A sync group must contain one cloud endpoint, which represents an Azure file share and one or more server endpoints.

A server endpoint represents a path on registered server.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-sync-files-deployment-guide>

HOTSPOT

You have an Azure subscription that contains the resources shown in the following table.

| Name | Type | Resource group |
|-------|-----------------|----------------|
| VNET1 | Virtual network | RG1 |
| VNET2 | Virtual network | RG2 |
| VM1 | Virtual machine | RG2 |

The status of VM1 is Running.

You assign an Azure policy as shown in the exhibit. (Click the Exhibit tab.)

Home > Policy - Assignments > Assign Policy

Assign Policy

SCOPE

* Scope (Learn more about setting the scope)
Azure Pass/RG2

Exclusions

Optionally select resources to exempt from the policy assignment

BASICS

* Policy definition
Not allowed resource types

* Assignment name ⓘ
Not allowed resource types

Description

Assigned by
First User

PARAMETERS

* Not allowed resource types ⓘ
3 selected

Assign Cancel

You assign the policy by using the following parameters:

Microsoft.ClassicNetwork/virtualNetworks

Microsoft.Network/virtualNetworks

Microsoft.Compute/virtualMachines

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

An administrator can move VNET1 to RG2 _____

Yes

No

The state of VM1 changed to deallocated _____

Yes

No

An administrator can modify the address space of VNET2 _____

Yes

No

DRAG DROP

You have an Azure subscription that contains a storage account.

You have an on-premises server named Server1 that runs Windows Server 2016. Server1 has 2 TB of data.

You need to transfer the data to the storage account by using the Azure Import/Export service.

In which order should you perform the actions? To answer, move all actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Select and Place:

1. From the Azure portal, update the import job
2. From the Azure portal, create an import job
3. Attach an external disk to Server1 and then run waimportexport.exe
4. Detach the external disks from Server1 and ship the disks to an Azure data center

3,2,4,1

1,4,3,2

1,2,3,4

2,3,4,1

Validate

**Solution:**

Explanation:

At a high level, an import job involves the following steps:

Step 1: Attach an external disk to Server1 and then run waimportexport.exe

Determine data to be imported, number of drives you need, destination blob location for your data in Azure storage.

Use the WAImportExport tool to copy data to disk drives. Encrypt the disk drives with BitLocker.

Step 2: From the Azure portal, create an import job.

Create an import job in your target storage account in Azure portal. Upload the drive journal files.

Step 3: Detach the external disks from Server1 and ship the disks to an Azure data center.

Provide the return address and carrier account number for shipping the drives back to you.

Ship the disk drives to the shipping address provided during job creation.

Step 4: From the Azure portal, update the import job

Update the delivery tracking number in the import job details and submit the import job.

The drives are received and processed at the Azure data center.

The drives are shipped using your carrier account to the return address provided in the import job.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-import-export-service>

You plan to use the Azure Import/Export service to copy files to a storage account.

Which two files should you create before you prepare the drives for the import job? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

an XML manifest file

a dataset CSV file

a JSON configuration file

a PowerShell PS1 file

a driveset CSV file

You have a Recovery Service vault that you use to test backups. The test backups contain two protected virtual machines.

You need to delete the Recovery Services vault.

What should you do first?

- From the Recovery Service vault, delete the backup data.
- Modify the disaster recovery properties of each virtual machine.
- Modify the locks of each virtual machine.
- From the Recovery Service vault, stop the backup of each backup item.

HOTSPOT

You have an Azure subscription named Subscription1 that contains the resources shown in the following table.

| Name | Type | Location | Resource group |
|----------|-------------------------|------------|----------------|
| RG1 | Resource group | West US | Not applicable |
| RG2 | Resource group | West US | Not applicable |
| Vault1 | Recovery Services vault | Central US | RG1 |
| Vault2 | Recovery Services vault | West US | RG2 |
| VM1 | Virtual Machine | Central US | RG2 |
| storage1 | Storage account | West US | RG1 |
| SQL1 | Azure SQL database | East US | RG2 |

In storage1, you create a blob container named blob1 and a file share named share1.

Which resources can be backed up to Vault1 and Vault2? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Can use Vault1 for backups: _____

- VM1 only
- VM1 and share1 only
- VM1 and SQL1 only
- VM1, storage1, and SQL1 only
- VM1, blob1, share1, and SQL1

Can use Vault2 for backups: _____

- storage1 only
- share1 only
- VM1 and share1 only
- blob1 and share 1 only
- storage1 and SQL1 only

Validate ✓

Solution:

Explanation:

Box 1: VM1 only

VM1 is in the same region as Vault1.

File1 is not in the same region as Vault1.

SQL is not in the same region as Vault1.

Blobs cannot be backup up to service vaults.

Note: To create a vault to protect virtual machines, the vault must be in the same region as the virtual machines.

Box 2: Share1 only.

Storage1 is in the same region (West USA) as Vault2. Share1 is in Storage1.

Note: After you select Backup, the Backup pane opens and prompts you to select a storage account from a list of discovered supported storage accounts. They're either associated with this vault or present in the same region as the vault,

but not yet associated to any Recovery Services vault.

Reference:

<https://docs.microsoft.com/bs-cyrl-ba/azure/backup/backup-create-rs-vault>

<https://docs.microsoft.com/en-us/azure/backup/backup-afs>

You have an Azure subscription named Subscription1.

You have 5 TB of data that you need to transfer to Subscription1.

You plan to use an Azure Import/Export job.

What can you use as the destination of the imported data?

- a virtual machine
- an Azure Cosmos DB database
- Azure File Storage
- the Azure File Sync Storage Sync Service

Validate



Solution:

Explanation:

Azure Import/Export service is used to securely import large amounts of data to Azure Blob storage and Azure Files by shipping disk drives to an Azure datacenter.

The maximum size of an Azure Files Resource of a file share is 5 TB.

Note:

There are several versions of this question in the exam. The question has two correct answers:

1. Azure File Storage
2. Azure Blob Storage

The question can have other incorrect answer options, including the following:

Azure Data Lake Store

Azure SQL Database

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-import-export-service>

HOTSPOT

You have an Azure subscription.

You create the Azure Storage account shown in the following exhibit.

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

The minimum number of copies of the storage account will be [answer choice] _____

- 1
- 2
- 3
- 4

To reduce the cost of infrequently accessed data in the storage account, you must modify the [answer choice] setting _____

- Access tier (default)
- Performance
- Account kind
- Replication

Validate ✓

Solution:

Explanation:

Box 1: 3

Locally Redundant Storage (LRS) provides highly durable and available storage within a single location (sub region). We maintain an equivalent of 3 copies (replicas) of your data within the primary location as described in our SOSP paper;

this ensures that we can recover from common failures (disk, node, rack) without impacting your storage account's availability and durability.

Box 2: Access tier

Change the access tier from Hot to Cool.

Note: Azure storage offers different access tiers, which allow you to store blob object data in the most cost-effective manner. The available access tiers include:

Hot - Optimized for storing data that is accessed frequently.

Cool - Optimized for storing data that is infrequently accessed and stored for at least 30 days.

Archive - Optimized for storing data that is rarely accessed and stored for at least 180 days with flexible latency requirements (on the order of hours).

Reference:<https://azure.microsoft.com/en-us/blog/data-series-introducing-locally-redundant-storage-for-windows-azure-storage/>

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers>

You have an Azure subscription named Subscription1.

You have 5 TB of data that you need to transfer to Subscription1.

You plan to use an Azure Import/Export job.

What can you use as the destination of the imported data?

an Azure Cosmos DB database

Azure Blob storage

Azure Data Lake Store

the Azure File Sync Storage Sync Service

You have an Azure subscription named Subscription1.

You have 5 TB of data that you need to transfer to Subscription1.

You plan to use an Azure Import/Export job.

What can you use as the destination of the imported data?

an Azure Cosmos DB database

Azure Blob storage

Azure Data Lake Store

the Azure File Sync Storage Sync Service

DRAG DROP

You have an Azure subscription that contains an Azure file share.

You have an on-premises server named Server1 that runs Windows Server 2016.

You plan to set up Azure File Sync between Server1 and the Azure file share.

You need to prepare the subscription for the planned Azure File Sync.

Which two actions should you perform in the Azure subscription? To answer, drag the appropriate actions to the correct targets. Each action may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place

1. Create a Storage Sync Service
2. Install the Azure File Sync agent
3. Create a sync group

Run Server Registration

1,2

2,3

4,1

3,4

HOTSPOT

You have an Azure subscription that contains the file shares shown in the following table.

| Name | Location |
|--------|----------|
| share1 | West US |
| share2 | West US |
| share3 | East US |

You have the on-premises file shares shown in the following table.

| Name | Server | Path |
|-------|---------|------------|
| data1 | Server1 | D:\Folder1 |
| data2 | Server2 | E:\Folder2 |
| data3 | Server3 | E\Folder2 |

You create an Azure file sync group named Sync1 and perform the following actions:

- Add share1 as the cloud endpoint for Sync1.
- Add data1 as a server endpoint for Sync1.
- Register Server1 and Server2 to Sync1.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

You can add share3 as an additional cloud endpoint for Sync1. _____

- Yes
 No

You can add data 2 as an additional server endpoint for Sync1. _____

- Yes
 No

You can add data3 as an additional server endpoint for Sync1. _____

- Yes
 No

Validate ✓

Solution:

Explanation:

Box 1: No

A sync group must contain one cloud endpoint, which represents an Azure file share and one or more server endpoints.

Box 2: Yes

Data2 is located on Server2 which is registered to Sync1.

Box 3: No

Data3 is located on Server3 which is not registered to Sync1.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-sync-files-deployment-guide?tabs=azure-portal%2Cproactive-portal#create-a-sync-group-and-a-cloud-endpoint>

You have an Azure subscription that contains the storage accounts shown in the following exhibit.

Storage accounts

| Default Directory | | | | | |
|--|---|---|---|---|---------|
|  Add |  Manage view |  Refresh |  Export to CSV |  Assign tags | |
|  Delete | |  Feedback | | | |
| Filter by name... | | Subscription == all  | Resource group == all  | Location == all  | |
| Showing 1 to 4 of 4 records | | | | | |
| <input type="checkbox"/> | Name ↑ | Type ↑ | Kind ↑ | Resource group ↑ Location ↑ | |
| <input type="checkbox"/> |  contoso101 | Storage account | StorageV2 | RG1 | East US |
| <input type="checkbox"/> |  contoso102 | Storage account | Storage | RG1 | East US |
| <input type="checkbox"/> |  contoso103 | Storage account | BlobStorage | RG1 | East US |
| <input type="checkbox"/> |  contoso104 | Storage account | FileStorage | RG1 | East US |

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

You can create a premium file share in _____

- contoso101 only
- contoso104 only
- contoso101 or contoso 104 only
- contoso101, contoso102, or contoso104 only
- contoso101, contoso102, contoso103, or contoso104

You can use the Archive access tier in _____

- contoso101 only
- contoso101 or contoso103 only
- contoso101, contoso 102 , and contoso 103 only
- contoso101, contoso102, and contoso104 only
- contoso101, contoso102, contoso103, and contoso104

Validate 

Solution:

Explanation:

Box 1: contoso104 only

Premium file shares are hosted in a special purpose storage account kind, called a FileStorage account.

Box 2: contoso101, contoso102, and contoso103 only

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-how-to-create-premium-fileshare?tabs=azure-portal>

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers>

You have two Azure virtual machines named VM1 and VM2. You have two Recovery Services vaults named RSV1 and RSV2.

VM2 is backed up to RSV1.

You need to back up VM2 to RSV2.

What should you do first?

- From the RSV1 blade, click Backup items and stop the VM2 backup
- From the RSV2 blade, click Backup. From the Backup blade, select the backup for the virtual machine, and then click Backup
- From the VM2 blade, click Disaster recovery, click Replication settings, and then select RSV2 as the Recovery Services vault
- From the RSV1 blade, click Backup Jobs and export the VM2 job

You have a general-purpose v1 Azure Storage account named storage1 that uses locally-redundant storage (LRS).

You need to ensure that the data in the storage account is protected if a zone fails. The solution must minimize costs and administrative effort.

What should you do first?

- Create a new storage account.
- Configure object replication rules.
- Upgrade the account to general-purpose v2.
- Modify the Replication setting of storage1.

You have an Azure subscription that contains the storage accounts shown in the following table.

| Name | Type | Performance |
|-----------|------------------|-------------|
| storage1 | StorageV2 | Standard |
| storage 2 | BlobStorage | Standard |
| storage3 | BlockBlobStorage | Premium |
| storage4 | File Storage | Premium |

You plan to manage the data stored in the accounts by using lifecycle management rules.

To which storage accounts can you apply lifecycle management rules?

- storage1 only
- storage1 and storage2 only
- storage3 and storage4 only
- storage1, storage2, and storage3 only
- storage1, storage2, storage3, and storage4

You have an Azure subscription.

In the Azure portal, you plan to create a storage account named storage1 that will have the following settings:

- Performance: Standard
- Replication: Zone-redundant storage (ZRS)
- Access tier (default): Cool
- Hierarchical namespace: Disabled

You need to ensure that you can set Account kind for storage1 to BlockBlobStorage.

Which setting should you modify first?

Performance

Replication

Access tier (default)

Hierarchical namespace

You create an Azure Storage account named contosostorage.

You plan to create a file share named data.

Users need to map a drive to the data file share from home computers that run Windows 10.

Which outbound port should you open between the home computers and the data file share?

80

443

445

3389

Validate



Solution:

Explanation:

Server Message Block (SMB) is used to connect to an Azure file share over the internet. The SMB protocol requires TCP port 445 to be open.

Incorrect Answers:

A: Port 80 is required for HTTP to a web server

B: Port 443 is required for HTTPS to a web server

D: Port 3389443 is required for Remote desktop protocol (RDP) connections

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-how-to-use-files-windows>

You have an Azure Storage account named storage1 that contains a blob container named container1.

You need to prevent new content added to container1 from being modified for one year.

What should you configure?

- the access tier
- an access policy
- the Access control (IAM) settings
- the access level

You have an Azure Storage account named storage1 that contains a blob container. The blob container has a default access tier of Hot. Storage1 contains a container named container1.

You create lifecycle management rules in storage1 as shown in the following table.

| Name | Rule scope | Blob type | Blob subtype | Rule block | Prefix match |
|-------|---------------------------------|------------|--------------|---|-----------------|
| Rule1 | Limit blobs by using filters. | Blockblobs | Baseblobs | If base blobs were not modified for two days, move to archive storage. If base blobs were not modified for nine days, delete the blob. | container1/Dep1 |
| Rule2 | Apply to all blobs in storage1. | Blockblobs | Baseblobs | If base blobs were not modified for three days, move to cool storage. If base blobs were not modified for nine days, move to archive storage. | not applicable |

You perform the actions shown in the following table.

| Date | Action |
|-----------|---|
| October 1 | Upload three files named Dep1File1.docx, File2.docx, and File3.docx to container 1. |
| October 2 | Edit Dep1File1.docx and File3.docx. |
| October 5 | Edit File2.docx. |

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

On October 10, you can read Dep1File1.docx. _____

Yes

No

On October 10, you can read File2.docx. _____

Yes

No

On October 10, you can read File3.docx. _____

Yes

No

Note: This question-is part of a series of questions that present the same scenario. Each question-in the series contains a unique solution that might meet the stated goals. Some question-sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question-in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription named Subscription1. Subscription1 contains a resource group named RG1. RG1 contains resources that were deployed by using templates.

You need to view the date and time when the resources were created in RG1.

Solution: From the Subscriptions blade, you select the subscription, and then click Programmatic deployment.

Does this meet the goal?

Yes

No

Note: This question-is part of a series of questions that present the same scenario. Each question-in the series contains a unique solution that might meet the stated goals. Some question-sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question-in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription named Subscription1. Subscription1 contains a resource group named RG1. RG1 contains resources that were deployed by using templates.

You need to view the date and time when the resources were created in RG1.

Solution: From the RG1 blade, you click Automation script.

Does this meet the goal?

Yes

No

Note: This question-is part of a series of questions that present the same scenario. Each question-in the series contains a unique solution that might meet the stated goals. Some question-sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question-in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure subscription named Subscription1. Subscription1 contains a resource group named RG1. RG1 contains resources that were deployed by using templates.

You need to view the date and time when the resources were created in RG1.

Solution: From the RG1 blade, you click Deployments.

Does this meet the goal?



Yes



No

You have an Azure subscription named Subscription1.

You deploy a Linux virtual machine named VM1 to Subscription1.

You need to monitor the metrics and the logs of VM1.

What should you use?

- Azure HDInsight
- Linux Diagnostic Extension (LAD) 3.0
- the AzurePerformanceDiagnostics extension
- Azure Analysis Services

You plan to deploy three Azure virtual machines named VM1, VM2, and VM3. The virtual machines will host a web app named App1.

You need to ensure that at least two virtual machines are available if a single Azure datacenter becomes unavailable.

What should you deploy?

all three virtual machines in a single Availability Zone

all virtual machines in a single Availability Set

each virtual machine in a separate Availability Zone

each virtual machine in a separate Availability Set

You have an Azure virtual machine named VM1 that runs Windows Server 2019.

You save VM1 as a template named Template1 to the Azure Resource Manager library.

You plan to deploy a virtual machine named VM2 from Template1.

What can you configure during the deployment of VM2?

operating system

administrator username

virtual machine size

resource group

You have an Azure subscription that contains an Azure virtual machine named VM1. VM1 runs a financial reporting app named App1 that does not support multiple active instances.

At the end of each month, CPU usage for VM1 peaks when App1 runs.

You need to create a scheduled runbook to increase the processor performance of VM1 at the end of each month.

What task should you include in the runbook?

- Add the Azure Performance Diagnostics agent to VM1.
- Modify the VM size property of VM1.
- Add VM1 to a scale set.
- Increase the vCPU quota for the subscription.
- Add a Desired State Configuration (DSC) extension to VM1.

You have an Azure virtual machine named VM1 that runs Windows Server 2019 . The VM was deployed using default drive settings.

You sign in to VM1 as a user named User1 and perform the following actions:

- Create files on drive C.
- Create files on drive D.
- Modify the screen saver timeout.
- Change the desktop background.

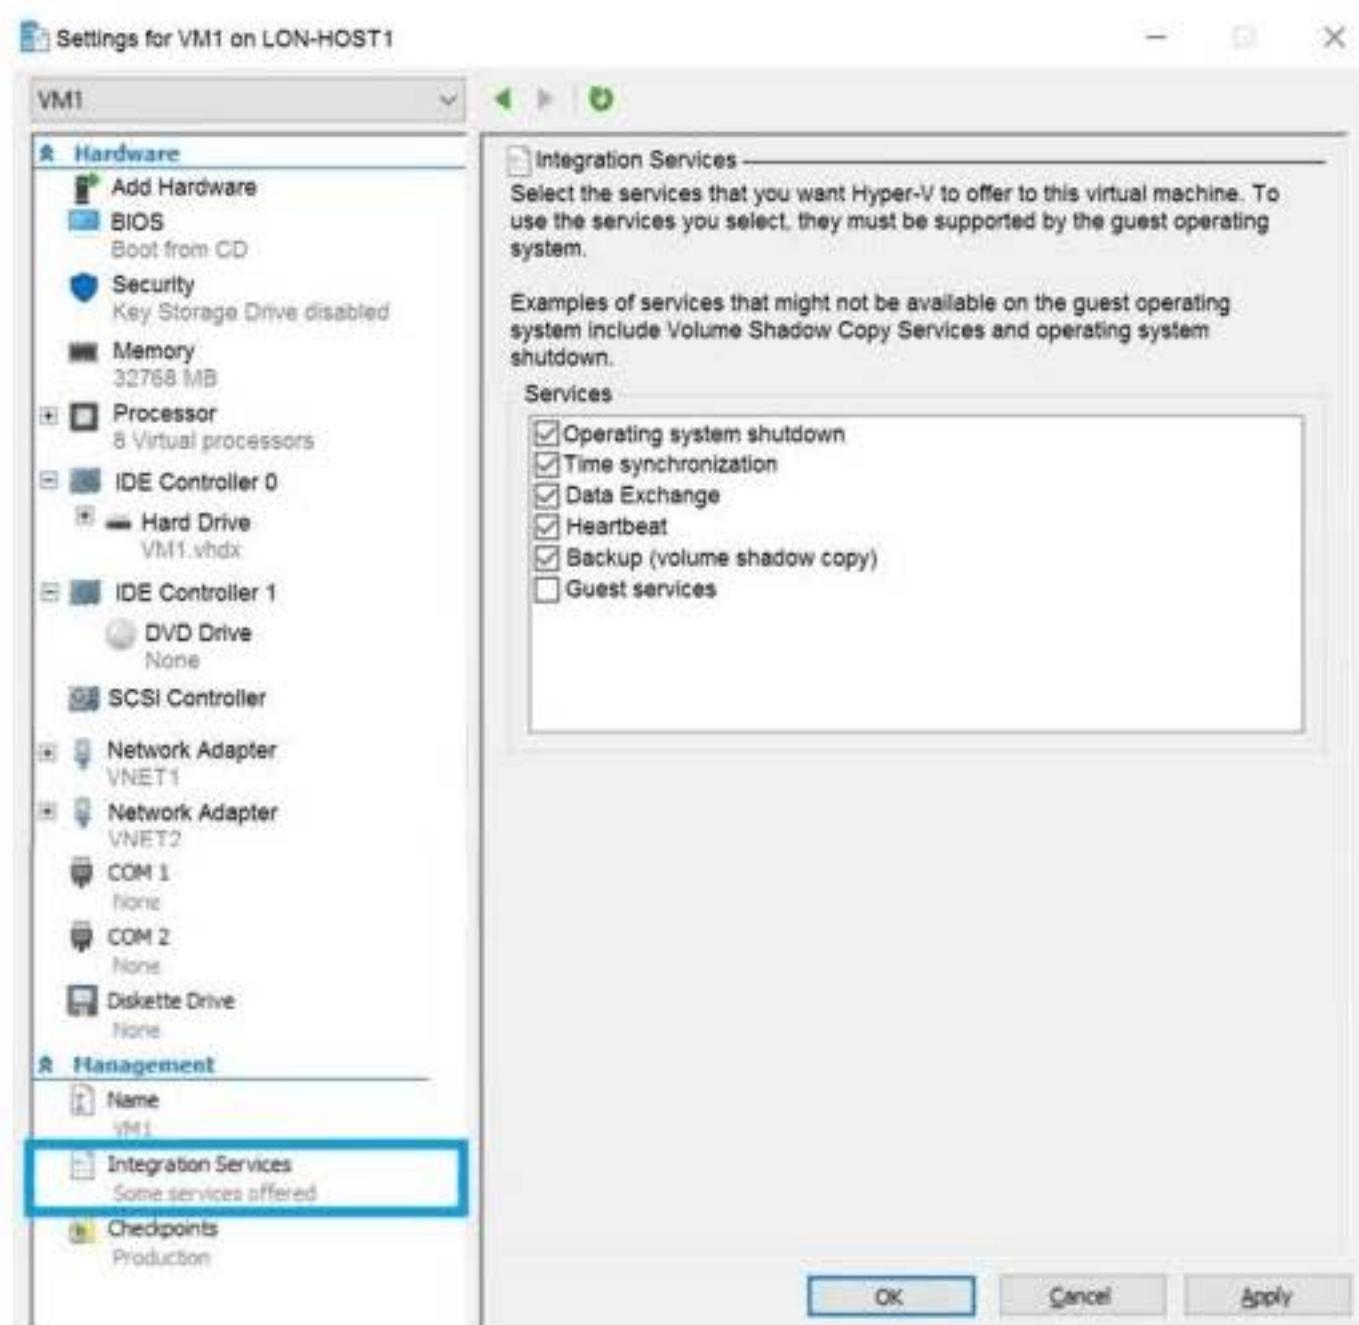
You plan to redeploy VM1.

Which changes will be lost after you redeploy VM1?

- the modified screen saver timeout
- the new desktop background
- the new files on drive D
- the new files on drive C

You have an Azure subscription.

You have an on-premises virtual machine named VM1. The settings for VM1 are shown in the exhibit. (Click the Exhibit tab.)



You need to ensure that you can use the disks attached to VM1 as a template for Azure virtual machines.

What should you modify on VM1?

- the memory
- the network adapters
- the hard drive
- the processor
- Integration Services

Validate ✓

Solution:

From the exhibit we see that the disk is in the VHDX format.

Before you upload a Windows virtual machine (VM) from on-premises to Microsoft Azure, you must prepare the virtual hard disk (VHD or VHDX). Azure supports only generation 1 VMs that are in the VHD file format and have a fixed size.

disk. The maximum size allowed for the VHD is 1,023 GB. You can convert a generation 1 VM from the VHDX file system to VHD and from a dynamically expanding disk to fixed-sized.

Reference:

<https://docs.microsoft.com/en-us/azure/virtual-machines/windows/prepare-for-upload-vhd-image>